



【ISSTA 24】 DBStorm: Generating Various Effective Workloads for Testing Isolation Levels

2024.09.24



中國人民大學
RENMIN UNIVERSITY OF CHINA

目录

CONTENTS

1

Background

2

DBStorm Framework

3

Experiments



中國人民大學
RENMIN UNIVERSITY OF CHINA

Part 1

Background



Background

- 冲突图(Conflict Graph)来表示数据库事务之间的并发冲突

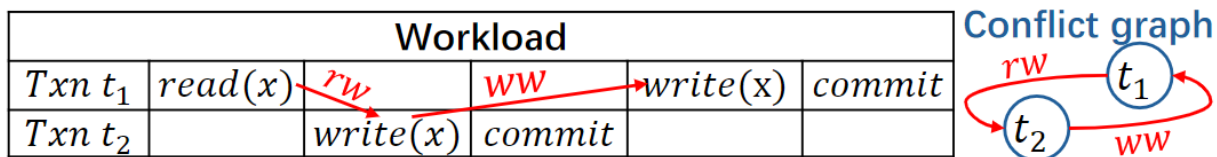


Figure 1: Example of the Conflict Graph

- 不同的异常对应特定的冲突环模式(Conflict Cycle Pattern)
- 现有的研究未能有效生成用于测试隔离级别的多样化负载，难以全面覆盖

Anomaly	Conflict Cycle Pattern	Likely to Occur In
Write-Skew	A cycle contains at least two consecutive rw conflicts	Read Committed, Repeatable Read,
Lost-Update	A cycle contains one rw conflict consecutive with one ww conflict	Read Committed, Repeatable Read
Read-Skew	A cycle contains one rw conflict consecutive with one wr conflict	Read Committed, Repeatable Read
Dirty-Read	A cycle contains at least one wr conflict but without rw conflict	Read Uncommitted
Dirty-Write	A cycle contains ww conflict only	Read Uncommitted



Challenges

C1: Black-box Testing : internal **I**solat**I**on **L**evel(**IL**) code is inaccessible

Generate a workload with a large number of conflict graphs.

C2: Effective Testing : **redundant** or **ineffective** workloads

A lightweight data state mirroring method: precise record access and efficiently generate dissimilar conflict graphs

C3: Anomaly-sensitive Testing : customize workloads to trigger **specific anomalies**

An implantation-based approach that orchestrates conflict record accesses

高效地生成精准且能够检测不同异常的workloads



中國人民大學
RENMIN UNIVERSITY OF CHINA

Part 2

DBStorm Framework



DBStorm Framework

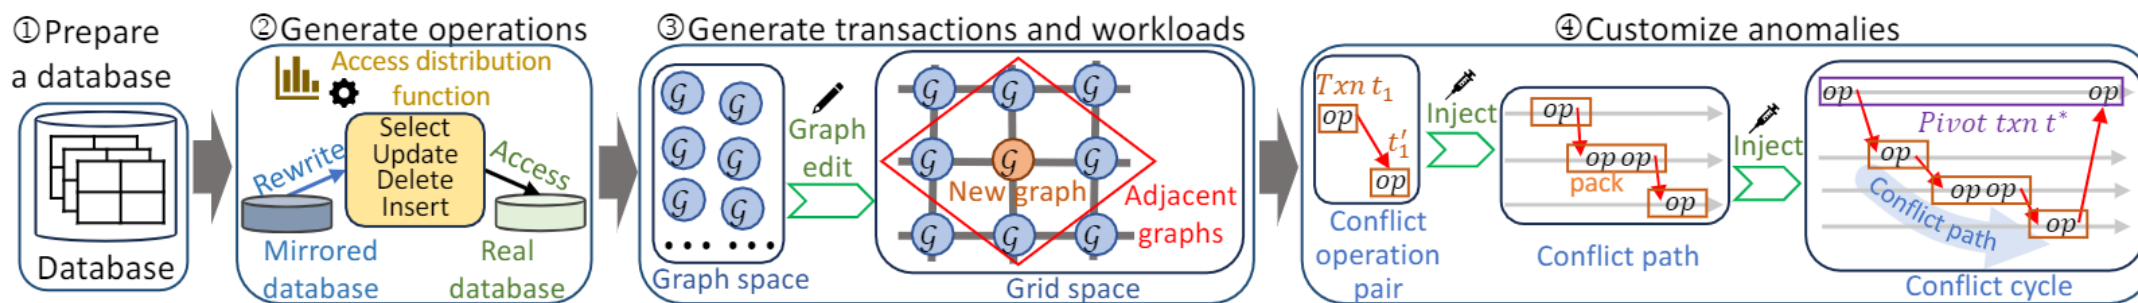


Figure 2: DBStorm Framework

- ① Preparing a DB instance: SQLancer/DT²
- ② Generating SQL Operations: Mirror the Data State and Instantiate Access Parameters
- ③ Grouping SQL Operations as Transactions: Conflict Graph Isomorphism Mapping
- ④ Customizing Non-serializable Anomalies: An implantation-based approach and a pivot based transaction scheduling policy.



Operation Generation

Challenge: ensure that each operation actually accesses its expected records

Mirroring the Data State

a distribution-guided method for mirroring and manipulating the fk/nk attributes.

R1-Insert: the data to be inserted is sampled from the data distribution function;

R2-Delete: the data to be deleted is randomly selected from the database;

R3-Update: performed by a deletion and then an insertion on the same record.

a partition-based method for mirroring and manipulating the pk attributes.

Static keys; Dynamic keys $\langle L, \lambda \rangle$, $L \in \{ \text{free}, \text{read}, \text{write} \}$, $\lambda \in \{0,1\}$;


Instantiating Access Parameters:

Item-read/write with pk

pk dynamic & $\lambda = 0$, insert; $\lambda = 1$, select update, delete; Static: update fk/nk

Predicate-read with fk/nk

$pred \leftarrow$ instantiate aps by sampling from the mirrored attribute;





Transaction Generation

Challenge: 生成足够多样化的冲突图，图的相似度通过Graph Edit Distance衡量

C1: 事务中的操作数量


Randomly select a transaction from the shrinking group and merge its operation into a growing group transaction.

C2: 事务读写操作比例


Randomly select a write operation and mutate it as a read operation.

history-independent & monotonically

Hybrid Editing Method: considers both C1's and C2's editing distances.



(a) Separate Editing Method



(b) Hybrid Editing Method



Anomaly Customization

An implantation-based approach

- ① 确定冲突环的模式 ϕ 和长度 L , 相邻两个操作(op_{2i} 和 op_{2i+1})需要访问相同的记录, 并且至少有一个是写操作
- ② 启动一个长事务(pivot事务)并注入操作 op_0
- ③ 将队列中的操作植入其他事务
- ④ 完成冲突环的构造

$\{op_0(R1), op_1(W1), op_2(R2), op_3(W2), op_4(R3), op_5(W3)\}$

$T1(pivot)$
 op_0-R1

$T2$

$T3$

op_1-W1
 op_2-R2

op_3-W2
 op_4-R3

op_5-W3




Anomaly Customization

An implantation-based approach

- ① 确定冲突环的模式 ϕ 和长度 L
- ② 启动一个长事务(pivot事务)并植入操作 op_0
- ③ 将队列中的操作植入其他事务
- ④ 完成冲突环的构造

$\{op_0(R1), op_1(W1), op_2(R2), op_3(W2), op_4(R3), op_5(W3)\}$





中國人民大學
RENMIN UNIVERSITY OF CHINA

Part 3

Experiments



IL-related Code Coverage

Baseline: [SQLsmith、Squirrel、SSQLancer、TQS] [TPC-C、SmallBank、YCSB]
[Jepsen、DT2、Troc]

A higher code coverage implies a higher likelihood of detecting bugs

Access Ratio α and Distribution Fitness β

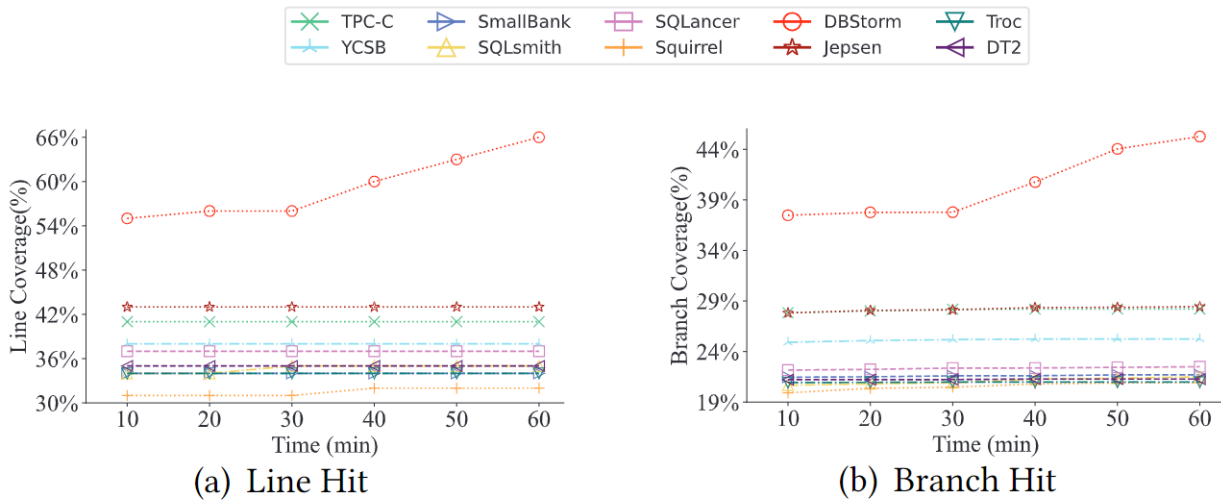


Figure 6: Code Coverage Comparison

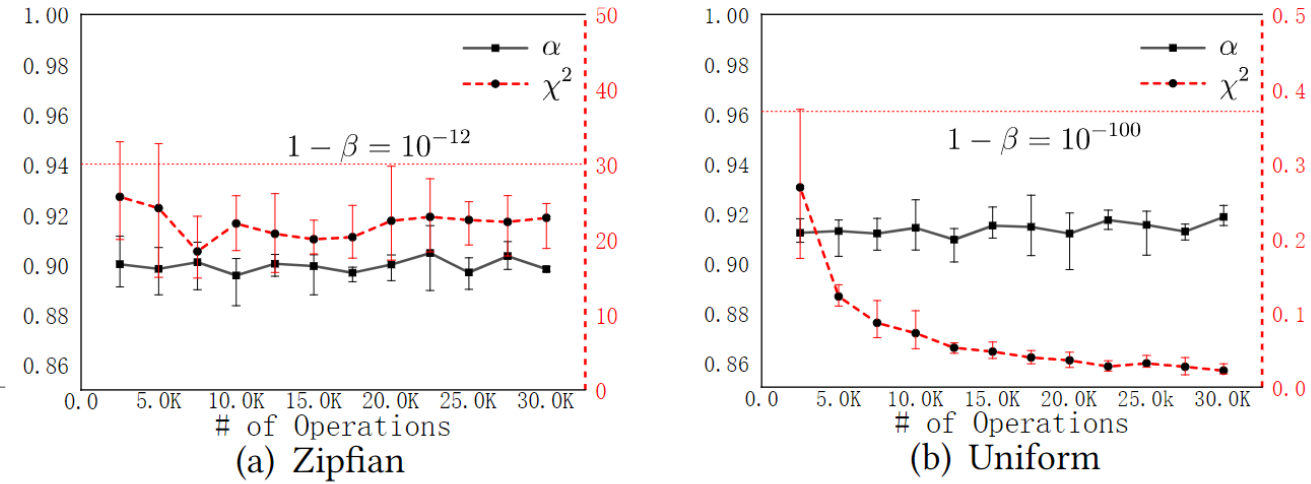


Figure 7: Access Ratio α and Distribution Fitness β



Effective Testing

Diversity of Conflict Graphs.

TQS: Embedding-based

$3\times$ higher than TQS in graph generation throughput and remains linear scalable

$2\times$ higher than TQS in generating non-isomorphic graphs

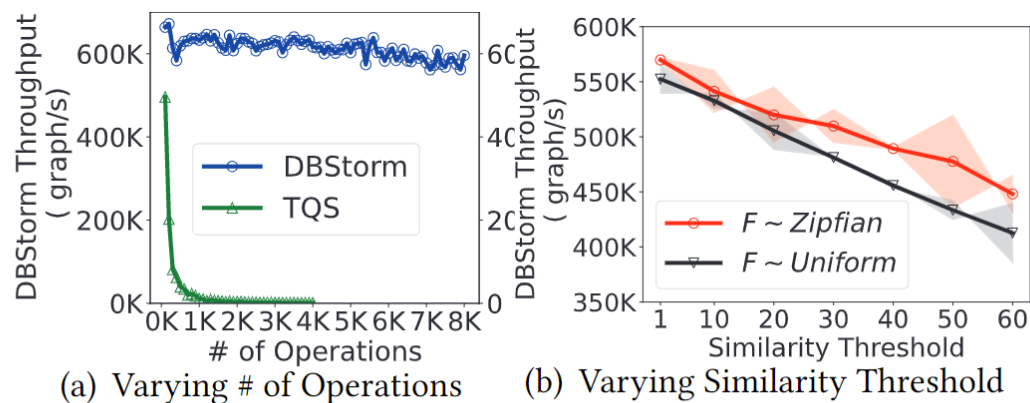


Figure 8: The Throughput of Diverse Graph Generation

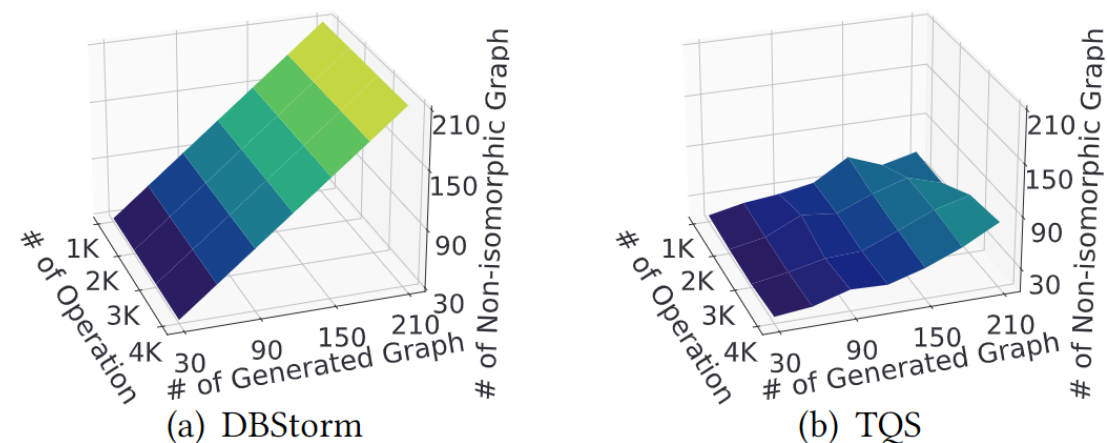


Figure 9: The Accuracy of Non-isomorphic Graph Generation



Anomaly-Sensitive Testing

Disable CC : DBStorm can form long conflict cycles and generate a wide variety of anomalies.

Enable CC : DBStorm can trigger more prevention actions than that of other tools

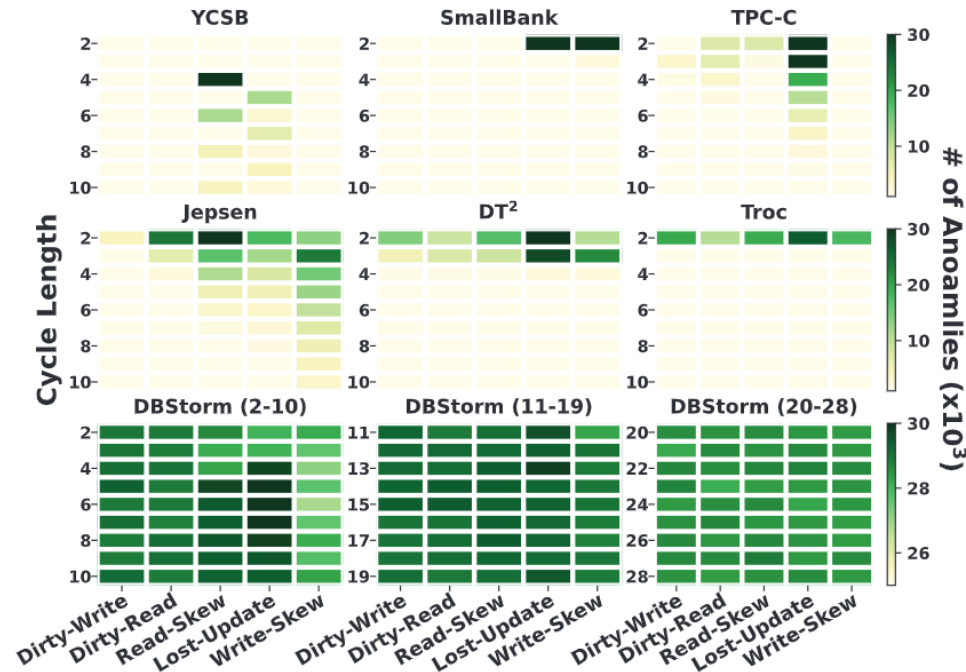


Figure 10: Anomaly Formation Comparison




Figure 11: Anomaly Prevention Action Comparison

Table 2: Bug Comparison

Total/Confirmed/Fixed	DBStorm	TxCheck	Jepsen	DT ²
Isolation	20/12/6	7/6/0	22/21/13	4/2/1
Non-isolation	13/12/10	49/46/18	33/33/20	12/10/0
Total/Confirmed/Fixed	Troc	SQLsmith	SQLancer	Squirrel
Isolation	5/5/1	0/0/0	0/0/0	0/0/0
Non-isolation	7/7/0	248/236/183	449/445/415	63/63/52